

<Heap Allocation>

1. Static vs. Local Variables

local variable { stored on the call stack: (push & pop)

Lifetime: function: Called return

params, return address, local vars are pushed by compiler

Static variable { stored at a fixed memory location (heap)

assigned at compile time by linker (loc is known)
encoded in the executable

Lifetime = (static) the complete runtime of the program

can be referenced from local variables

cons: read-only by default → should be encapsulated in a 'Mutex' type

2. Dynamic Memory = Heap

(a single mutex at any point)

- heap supports dynamic memory allocation

⇒ avoid data race

at runtime through [allocate] & [deallocate].

- Common errors
- ① memory leak: forget to deallocate (excessive mem consumption)
 - ② use-after-free: can be exploited to exe arbitrary code
 - ③ double-free: can lead to use-after-free.

⇓

Garbage Collection: the program is regularly paused & scanned for (Java, Python...) unused heap variables ⇒ auto deallocate

⇓ (assign an abstract lifetime to each reference)

Rust: Ownership: - check the correctness of dynamic memory operations at compile time

alloc: Box::new(...)

dealloc: Drop trait

(goes out of scope) ⇒ no performance overhead

⇒ no garbage collection at runtime & fine-grained ctrl

Resource acquisition is initialization (RAII)

Rust ensures memory safety & thread safety

<Allocator Design>.


① to manage the available heap memory.

- goals $\left\{ \begin{array}{l} \Rightarrow \text{return unused memory on 'alloc'} \\ \text{keep track of memory freed by 'dealloc' so that it can be reused.} \\ \text{effectively utilize the available memory \& keep fragmentation low.} \\ \text{work well for concurrent apps. \& scale to any \# processors.} \\ \text{optimize the memory layout w.r.t. CPU caches to \uparrow cache locality. \& avoid false sharing.} \end{array} \right.$

1. Bump Allocator (stack allocator)

- allocates memory linearly \& only keeps track of \# allocated bytes \& \# allocations.

- Design idea:



It is often implemented with an allocation counter:

- +1 on each 'alloc' call
- 1 on each 'dealloc' call

when [counter] == 0, it means that all allocations on the heap have been deallocated.

\Rightarrow [next] will be reset to the [heap-start]

\Rightarrow bump allocator can only free all memory at once!

pros: very fast - can be optimized to just a few assembly instrs.

cons: it can only reuse deallocated memory after all allocations have been freed.


\downarrow
[test-case]: many-boxes - long-lived.

① only splits the heap into smaller blocks but never merges them back.
(ext fragmentation)

com } ② worse performance

2 Linked List Allocator

- constructing a single linked list in the freed memory to store the information about freed region.






each list node (free list) / the size of the memory region or pointer to the next unused memory region.

2.1. Implementation of [add-free-region]

takes the 'addr' & 'size' of a mem region,
adds it to the front of the list.

① checks that the given region has the necessary size & alignment for storing a 'ListNode'.

② creates the node & inserts it into the head of the list



2.2. [find-region]: iterate over the free list to find suitable region.

2.3. [alloc-from-region]: given a 'region', 'size', 'align'

① calculate the start & end address of a potential allocation

(alloc-start = align-up(region.start, align))

② calculate the excessive size \rightarrow $\begin{cases} > 0 \&\& \geq \text{sizeof(ListNode)} \\ = 0 \end{cases}$

③ return the 'alloc-start'

3. Fixed-Size Block Allocator.

- idea: Instead of allocating exactly as much memory as required,

we define a small # of block sizes & round up each allocation to the

Also keep track of the unused mem by creating a linked list ^{next block size} in the unused memory.

But create a separate list for each size class.

(e.g. head-16, head-64, head-512)

- alloc: ① Choose the smallest mem block that satisfies the request.

② retrieve the head pointer for the list (e.g. head-16)

③ remove the first block & return it

- dealloc: ① round up the freed allocation size to the next block size.

② retrieve the head pointer for the list

③ add the freed block to the front of the list & update head ptr.

- create new blocks: allocate new block from fallback allocator

(2 ways): split a larger block from a different list.

- Fallback allocator:

Large allocations (> 2KB) are often rare, especially in OS kernels.

It might make sense to fallback to a different allocator for these allo.

Pros: performance

Cons: internal fragmentation

- Variations:

Slab allocator

Buddy allocator